



- random init (78.1%)
- proc weights in all blocks (80.1%)
- blocks 0-8: proc weights, blocks 9-11 random (80.0%)
- blocks 0-8: per-matrix weight std from proc (zero mean), LN gains 0.3-0.5 (79.9%)

- blocks 9-11: proc weights, blocks 0-8 random (1 seed) (78.9%)
- blocks 9-11: v, proj, fc2 scaled to $\|f(x)\|/\|x\| = 1.4$ (79.7%)
- blocks 9-11: v, proj, fc2 scaled to proc's $\|f(x)\|/\|x\|$ (80.0%)


(a) Residual-write ratio at initialisation



(b) Class decodability per block, epoch 289



(c) Probe accuracy during training



(d) Residual-write ratio at epoch 289

