

The gain is regularisation: winners fit the clean training images less and generalise better

Final checkpoints, 25k training images without augmentation and 25k validation images. Spearman(gap, accuracy) = -0.94; clean-train top-1 is 99.3-99.4% for random-level and losers, 96.6-98.4% for prefix-like winners.

