

Random-init early blocks scramble tokens from step 0; every winner starts transparent and opens over 10-30 epochs

Blocks 1-7, 256 fixed training images, per-epoch checkpoints. On ksd the MLPs wake within ~2 epochs and only the token-specific attention stays low for ~10.

